

Shorthand Note: Overarching Formalism & Operationalization of the Multisensor Fusion Lake

1 Overarching Formalism

Metadata–Signal Factorisation. Every geospatial measurement is treated as a pair

$$\left(\underbrace{m}_{\text{metadata}}, \underbrace{s}_{\text{signal}} \right) \in \text{Meta} \times \text{Signal},$$

where **Meta** is a STAC¹ object and **Signal** is a Zarr store. The product structure guarantees that changes in calibration, reprojection grids, or statistical models never mutate the physical bytes of other factors.

Sensor Alignment. Let $S_i : \mathcal{G} \rightarrow \mathbb{R}^{k_i}$ be the field generated by the i -th sensor after applying its affine reprojection T_i^{-1} . With m sensors the aligned bundle is

$$\mathbf{S}(\mathbf{x}) := (S_1(\mathbf{x}), \dots, S_m(\mathbf{x})), \quad \mathbf{x} \in \mathcal{G} \subset \mathbb{R}^3.$$

Fusion Functional. For a task-specific function $g : \mathbb{R}^{k_1 + \dots + k_m} \rightarrow \mathbb{R}^d$ we estimate the posterior mean

$$F(\mathbf{x}) = \mathbb{E}[g(\mathbf{S}(\mathbf{x})) \mid \text{data, hyper-priors}],$$

evaluated chunk-wise so that computation factorises over Zarr keys.

Reproducibility Ledger. Each execution appends

$$\Gamma := [t, \text{item_hash}, \text{posterior_hash}, \text{ELBO}, E, \tau]$$

to an append-only log; cryptographic hashes make the entire pipeline provably referentially transparent.

2 Operationalization

2.1 Physical Layout (S3 / GCS / MinIO)

```
s3://fusion-lake/
  raw/           # original sensor blobs
  stac/          # catalog, collections, items
  posterior/     # derived Zarr groups (versioned)
```

¹SpatioTemporal Asset Catalog.

2.2 Ingestion DAG

1. **Stage**: wrap new object in a STAC `Item`; attach checksums.
2. **Reproject**: convert native grid \rightarrow common \mathcal{G} ; write `asset.type = application/vnd+zarr`.
3. **Posterior Notebook**: parameterised execution (Papermill) computes (μ, σ^2) or other sufficient statistics; emits new Zarr under `posterior/`.
4. **Ledger Write**: hash `Item` and Zarr root; append to Γ .

2.3 STAC \leftrightarrow Zarr Contract

- Every STAC asset whose `type` is `application/vnd+zarr` *must* resolve to a store whose root contains the JSON key `"zarr_format"`.
- A STAC extension field `"xarray:open_kwargs"` specifies the loader's arguments, enabling pure-metadata discovery by Jupyter / LangChain agents.

2.4 Notebook Template (Skeleton)

```
cat = intake.open_stac_catalog("s3://fusion-lake/stac/catalog.json")
item = cat["sentinel-2"][ITEM_ID]
ds    = item["B04_zarr"].to_dask()    # lazy xarray
# --- science here -----
posterior = bayes_update(ds)
posterior.to_zarr("s3://fusion-lake/posterior/ndvi_v0.zarr", mode="w")
```

2.5 Key Guarantees

Immutability All STAC Items are write-once; updates create new Item IDs.

Chunk Locality Fixed 512×512 window $\Rightarrow O(1)$ object-store reads.

Schema Versioning Evolution of Meta or Signal is tracked by semantic version tags in `stac/collections`.